


Deep-Q Learning Tetris AI

Ijahmin Balser



Problem

- Design an AI to score high in Tetris
- Tetris has a lot of possible states
 - Representation of game state can make this number even larger
- Smaller state space will lead to faster training and better performance for AI
- Common Tetris strategies tend towards long term pay-offs
 - Well designed reward function can help with short-term pay-offs





Use-Case Scenarios

- Reference for new Tetris players
- AI opponent in multiplayer Tetris games
- DQN Agent is transferrable other games/applications
 - Bug testing possibilities
 - Performance evaluation


AI Algorithm and Model

- Reinforcement Learning
 - Standard Q-Learning Algorithm
 - Deep Q-Learning Agent
- Game state evaluated based a few factors
 - Lines cleared
 - Number of holes
 - Bumpiness
 - Total height




Results


- Performance is heavily dependent on reward function
- Q-Learning policy produces little change when compared to a random policy
- Deep Q-Learning Agent improved rapidly
 - Episodes length becomes too long to record towards end of training set




Demo



Episode 510



Episode 1180



Episode 1630

Lessons Learned

- Minimizing representation of game leads to a much easier time training
- Relying on exclusively long-term payoffs is not reliable
- Optimized reward function is one of the most important factors

References

- <https://codemyroad.wordpress.com/2013/04/14/tetris-ai-the-near-perfect-player/>
- <https://github.com/nuno-faria/tetris-ai>
- <https://pytorch.org/docs/stable/tensorboard.html>
- <https://keras.io>