

Boop - AI Agent

Alexander Barry, Scott Berlow, Christopher Mascis

Problem Statement and Analysis

- Design an AI Agent to play *boop*
- Evaluate several AI algorithms
- Create a user interface to challenge the AI
- Implement using Python



Smirk & Dagger Games


Characteristics of *boop*

- Two-player
- Zero-sum
- Fully Observable
- Turn-taking
- Multiple Win Conditions



General Gameplay

Booping kittens




Graduating kittens



From the *boop* Rulebook

Additional Rules


- Multiple "triples" in one turn
 - Player gets to decide which kittens to promote
- Triples for both players in one turn
 - Each player gets to promote
- All eight pieces on the board
 - Player chooses one to promote



From the *boop* Rulebook

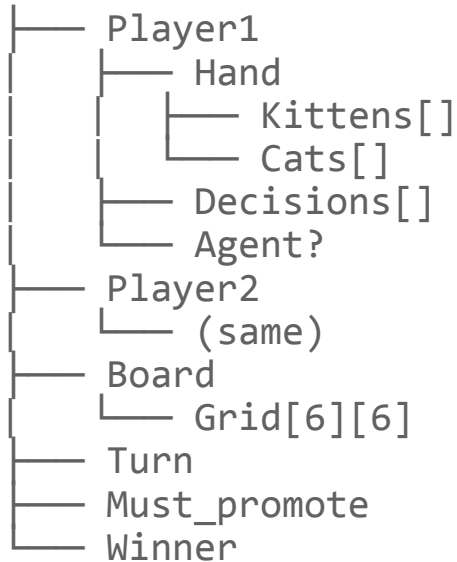
Use-Case Scenarios

- Human vs Human:
 - play the board game in a virtual UI against another human
- Human vs AI:
 - play without another person present
 - practice against an AI to hone skills
- AI vs AI:
 - self-play mode
 - evaluate agent types against each other
 - evaluate heuristic functions and potential features
 - explore strategies



State Representation

GameState



MinimaxAgent picked (K(2), 2, 3) with value
-2.65, states explored: 12097

AI action: K(2) to (2, 3)

0						K(2)
1	K(1)					
2				K(2)		
3						
4		K(1)				
5						
	0	1	2	3	4	5

AI Algorithm and Model

- Minimax
 - Configurable depth (where depth accounts for a layer of MAX and MIN each)
 - **Problem: too many states!**
 - Worst case: $\sim 36^{(d*2)}$ possible *spaces* to consider
 - If a player has cats and kittens, 2 possible moves *per space* at each level
 - Initial move with depth 2: $\sim 36^4 = 1,679,616$ states
 - With cats: $\sim 72^4 = 26,873,856$ states
 - Configurable `maxStates` parameter to cap exploration - but leads to suboptimal moves
- Alpha-Beta Pruning
 - Greatly reduces overall computation, though still can be slow
- Beam Search
 - Configurable beam "width" parameter
 - Forward pruning: at each level, maintain only the best n states, where n is the beam width
 - Shannon's Type B strategy: narrow but deep search

Evaluation Functions

- `eval_piece_count`
 - number of kittens and cats
 - cats are worth more
 - simple, but not very useful
- `eval_board_bonus`
 - pieces on board are more valuable than in hand
- `eval_territory`
 - center area more valuable
 - 2-in-a-row bonus
 - 7 on board bonus
 - L-shapes bonus
- `eval_stranding`
 - extra penalty if cats isolated on board, with no more in hand

```
def eval_territory(id: PlayerID, state: GameState):
    p1_score = _get_territory_score(PlayerID.ONE, state)
    p2_score = _get_territory_score(PlayerID.TWO, state)

    # apply anti-aggression measure
    if id == PlayerID.TWO:
        p1_score *= ANTI_AGGRESSION
    elif id == PlayerID.ONE:
        p2_score *= ANTI_AGGRESSION

    ply_penalty = _get_ply_penalty(id, state)
    win_bonus = _get_win_bonus(state)
    if win_bonus != 0:
        return win_bonus + ply_penalty

    return p1_score - p2_score + ply_penalty
```

Tuning Parameters

WIN_BONUS = 1000.0

PLY_PENALTY = 0.1

ANTI_AGGRESSION = 0.8

a cat must be worth more than 3x kittens

CAT_MULTIPLIER = 25.0

KITTEN_MULTIPLIER = 1.0

BOARD_MULTIPLIER = 2.0

CENTER_MULTIPLIER = 2.5

PENDING_TRIPLE_BONUS = 4.0 # encourage getting triples

PENDING_PROMOTION_BONUS = 3.5 # encourage getting all pieces onto board

STRANDED_CAT_PENALTY = 1.0

Results

Evaluation Functions - *conducted with AlphaBetaAgent, maxStates 500,000 and depth 2*

1. eval_piece_count vs **eval_board_bonus** - 56 plies
2. **eval_board_bonus** vs eval_territory - 55 plies
3. **eval_board_bonus** vs eval_stranding - 55 plies
4. **eval_territory** vs eval_board_bonus - 77 plies

Beam Parameters - *eval_territory, ANTI_AGGRESSION changes to prevent deadlock*

1. **10 width, 10 depth** vs 5 width, 25 depth - 47 plies, ANTI_AGG = 0.5
2. **20 width, 5 depth** vs 10 width, 25 depth - 53 plies, ANTI_AGG = 0.3
3. **10 width, 25 depth** vs 20 width, 5 depth - 49 plies, ANTI_AGG = 0.5

Alpha-Beta vs Beam Search - *eval_territory used for both*

1. Beam Search 10 width, 25 depth vs **AlphaBeta depth 2** - 22 plies
2. **AlphaBeta depth 2** vs Beam Search 10 width, 25 depth - 27 plies

Demonstration

Lessons Learned

- Evaluation functions are hard to get right!
 - Sometimes a small adjustment to one parameter makes a huge difference in behavior
 - AI can get stuck in cycles and deadlocks, especially vs itself
 - AI doesn't always go for the win, as it wants higher utility
- Rare cases in rules can take the longest to code
 - The "decisions" add complexity to minimax
 - This is especially true if players must act out of regular turn flow
- Visualization is a key for debugging
- Unit test the game logic - saves a lot of time and bugs in the long run

Future Work

- Better evaluation functions
- Optimization: can we exploit grid properties to save computations?
 - Example: equivalent moves based on grid rotation
- Reinforcement Learning Agent
 - Extract features from the Tuning Parameters and let the model train the weights
- Command line options - include in UI

Q/A