

Settlers of Catan AI

By: Jonathan West, Vincent DiPerna, Chiraag Potharaju

Problem Statement and Analysis

- Catan is a popular board game, and has a large user base
- There is a continual push to improve play
- Effective AI Game Playing Agents can help players improve
- This is shown with existing AI agents for other games



Catan Overview

- 2-4 Person Game
- Win by obtaining most “Victory Points”
- Build & expand settlements, harvest resources, play development cards, and trade with or rob other players
- Multiple Hexes (Resources),
 - Number on each hex matters
 - Positioning of hexes matters
- Multiple Actions per Turn (Order matters)



Use Cases

1

Catan Players

- Test new strategies
- Generate database of games to review
- Practice without needing other people
- Learn new strategies from AI

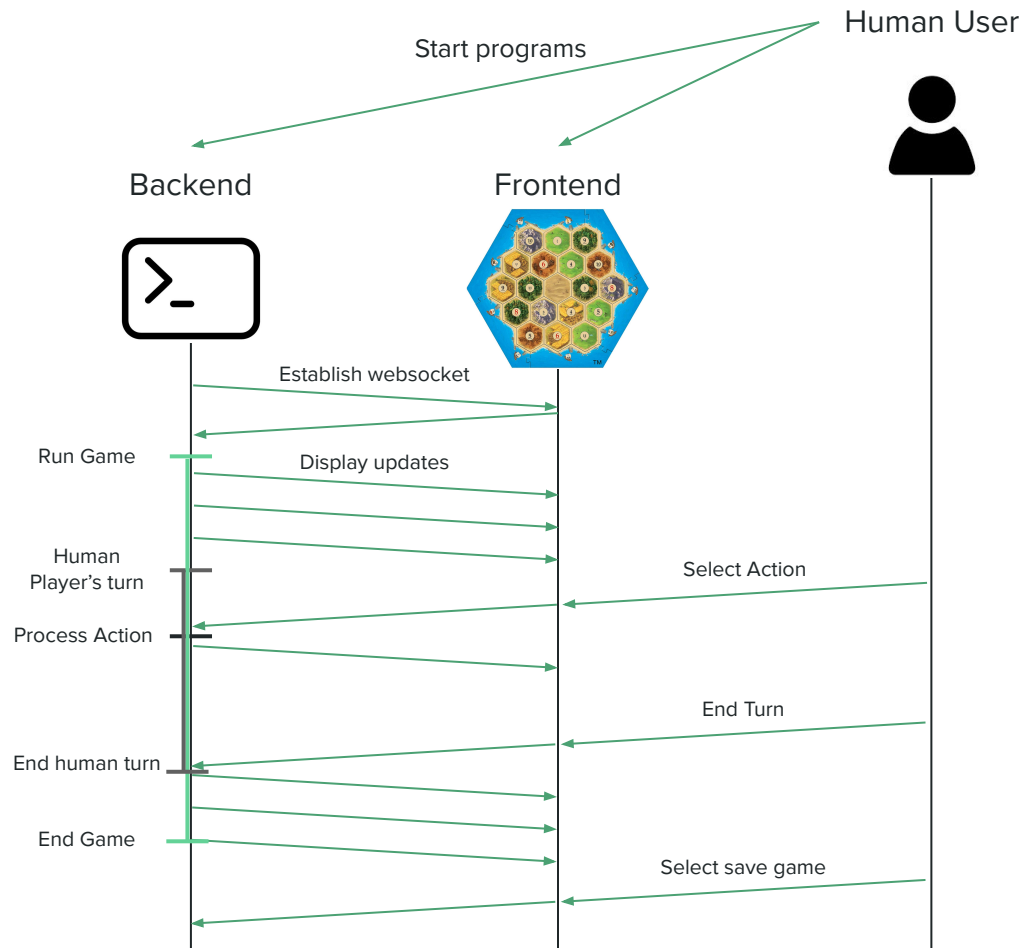
2

AI Researchers & Developers

- Generate large database of strong Catan games
- Apply similar AI models & techniques to new yet similar applications
- Learn from flaws in design when developing new projects

Basic Design

- Python backend game logic
 - Run in terminal with params for settings
- HTML Canvas frontend
 - Logic written with vanilla JavaScript
 - Communication via Websockets
 - Separated frontend logic to keep efficient backend
- AI Agents placed in additional files



AI Algorithm and Model: Considerations



Expectiminimax

- Allows for use with random chance & incomplete info, as opposed to minimax
- Primary AI agent



Monte Carlo Tree Search

- Handles uncertainty well
- More memory efficient than minimax
- Did not have time to implement



Deep Learning

- Could use a neural net as a state evaluation function
- Requires significant time & computational resources



Reinforcement Learning


- Catan AI learns from playing many instances of Catan under different circumstances

AI Algorithm: Expectiminimax

- Like with pacman ghosts, doesn't immediately alternate between two players
- Maximizes/minimizes with same player until they choose to end turn
- After a player ends turn, considers probabilistic dice roll before going to next player's turn in game tree
- Considered several game state features for evaluation heuristic, including player's VPs, opponents visible VPs, number resource & development cards held by each opponent
 - Balance of efficiency with depth of analysis

Design Considerations


- Balance between efficiency & human convenience
 - Could have used C rather than Python for algorithmic agent
 - Could have made somewhat more compact board state representation
 - Used hybrid bitmaps for board representation
- Flexibility for non-standard game experimentation
 - E.g. don't just assume there are always the same set of resource tiles in the board; allow for this variable to be easily changed
 - Balanced with pragmatic restrictions: e.g. didn't design program to easily allow multiple robbers
- Balancing memory spatial efficiency with computational efficiency
 - Example: Which players get which resources after a given dice roll?




Design Considerations

- Multi-stage decision making can reduce branching factor
 - Allows more specific heuristic implementations
- Generally leaned towards single-stage decision making
 - Used multiple decision stages for a player's turn, but not within a given building action
 - Simpler implementation

- Build city
- Build Settlement
- Build Road
- End Turn



Multiple Stages: Considered 8 actions



Single Stage: Considered 10 actions

Results

- Visual display similar to real board
- Allows natural human input using mouse
- Can generate large numbers of saved games for training, analysis, or other purposes
- Can save or load game states to and from .json files, and continue play simply
- Expectimax decent agent versus humans; still being improved upon
 - 100% win rate against random in our testing
- Relatively efficient game state in terms of common operation speed and size (~4000KB), while maintaining readability
- Several simplifying constraints for ease of development



```
7
10 # Road map, 1 bit for road locations
11 RMAP = [
12     0,
13     0b0000000000000000101010101000000,
14     0b0000000000000000100010001000000,
15     0b0000000000000000101010101010000,
16     0b0000000000001000100010001000100,
17     0b0000000000001010101010101010100,
18     0b0000000000100010001000100010010,
19     0b0000000000010101010101010101010,
20     0b0000000000001000100010001000100,
21     0b0000000000001010101010101010000,
22     0b000000000000000010001000100010000,
23     0b00000000000000101010101000000,
24     0]
25
```

Demonstration

Lessons Learned

- Furthered knowledge on heuristic design with relatively complex game
- Team organization & time management
- Would have implemented at least basic versions of a few more AI agents, but we lacked enough time
- Should have made more specific plan from the outset
- Optimize for common case
- Need to be careful with game simplifications: they can compound and lead to issues
- Should have spent more time on primary evaluation function; the current version was quite simplified for efficiency

Future Work

- Implement more AI agents, notably MCTS
- Continue accounting for parts of the game that were left out
- Improve heuristic evaluation function

Q&A

Try out the program:

<https://github.com/jrgranadoswest/ocatan>
